

Krishi Mitra — Integrated TODO List (Pipeline-Ordered)

Last updated: 2026-08-25 **Sprint goal:** Reach 9/10 prototype score in 2 weeks **Ordering principle:** Pipeline dependency first, then priority. Each task unblocks the next.

Legend

Symbol	Meaning
★	Demo-killer —fix first or examiner sees it fail
🔧	Infrastructure —enables other tasks
📊	Evaluation —needs working pipeline + gold labels
📝	Report writing —do last
NEW	NEW —not in original TODO.md
✓	DONE

Sprint Block 1: Foundation (Days 1–2)

These must be done first — everything else depends on them

#	Task	Effort	Status	Original TODO #	Notes
1.1	★ Test new Groq model (openai/gpt-oss-20b or verified active model)	1 hr	NEW NEW	#1	Smoke test: <code>python -c "from llm_client import call_llm; ..."</code>
1.2	★ Copy all 10 corrected files to backend/	30 min	NEW NEW	—	See CHANGES_INDEX.m for file list
1.3	🔧 Install deps: rank_bm25, openai, indic-nlp-library	15 min	NEW NEW	—	One command: <code>pip install rank_bm25 openai indic-nlp-library</code>

Table 2 – continued

#	Task	Effort	Status	Original TODO #	Notes
1.4	🔧 Set up back-end/.env from .env.example with real keys	15 min	 NEW	#6	Must have working GROQ_API_KEY + GENERATION_MODEL
1.5	🔧 Rebuild Qdrant database with semantic chunking + Indic NLP	10 min	 NEW	#13	python vector_db.py --stop main.py first
1.6	★ Re-label eval/gold.jsonl (chunk IDs changed after semantic chunking)	2–3 hrs	 NEW	—	Use chunk dump command from CHANGES_SUMMARY. ALL eval numbers are garbage until this is done.

Block 1 exit criteria: LLM responds, database rebuilt, gold labels are fresh.

Sprint Block 2: Demo-Killers (Days 2–5)

These fix the things examiners will notice in the first 2 minutes of your demo

#	Task	Effort	Status	Original TODO #	Notes
2.1	★ Query condensation for multi-turn retrieval	4-6 hrs	—	#14	BIGGEST DEMO RISK. Rewrite follow-up queries using chat history before embedding. Example: “ ?” → “ ?”
2.2	★ Streaming responses (SSE backend + frontend)	3-4 hrs	—	#15	Groq supports streaming. Frontend needs <code>EventSource</code> . Makes 30s feel like 5s.
2.3	★ Expand eval set to n=40-50 from Kisan Call Centre dataset	4-5 hrs	—	#18	n=16 is statistically thin. 40+ is defensible. Mine real Karnataka farmer questions. Example: “ ?” → “ ”
2.4	 Lightweight query expansion (LLM rewrites vague queries before retrieval)	2 hrs	 NEW	—	

Table 3 – continued

#	Task	Effort	Status	Original TODO #	Notes
2.5	 Adversarial test suite (numeric transposition, injection, empty query)	2 hrs	 NEW	—	Test: 10g→100g hallucination, off-topic injection, PII in query, empty string

Block 2 exit criteria: Follow-up questions work, streaming is live, eval set is 40+, adversarial tests pass.

Sprint Block 3: Evaluation & Metrics (Days 5–7)

These need Block 1 (fresh gold labels) + Block 2 (working pipeline)

#	Task	Effort	Status	Original TODO #	Notes
3.1	 Run threshold sweep with fresh gold labels	30 min	—	#7	python eval/thresh-old_sweep.py → pick best F1 threshold → update rag_service.py
3.2	 Run full retrieval ablation (dense/ sparse/ hybrid/ BM25/ bucketed)	30 min	—	#10	python eval/run_retrieval_a → save results for report
3.3	 Run generation eval with new model + cross-judge	1–2 hrs	—	#2, #4	python eval/run_eval.py with GENERATION_MODEL set. Use OpenRouter as judge if possible.

Table 4 – continued

#	Task	Effort	Status	Original TODO #	Notes
3.4	 Run refusal test on unanswerable questions	15 min	—	#11	python eval/re-fusal_test.py → document refusal rate
3.5	 Run numeric faithfulness demo	5 min	—	#5	python eval/numeric_faithfulness_demo.py → screenshot for report
3.6	 Run metric sanity check (chrF vs ROUGE-L vs embedding sim)	15 min	—	#3	python eval/diag-nose_metrics.py → confirm chrF is usable
3.7	 Manual failure-mode review (read bad answers, classify why)	2 hrs	—	#11	Read low-scoring answers. Categories: retrieval miss, generation hallucination, faithfulness false positive, etc.

Block 3 exit criteria: All eval scripts run, numbers are fresh, threshold is justified, failure modes are classified.

Sprint Block 4: Infrastructure Polish (Days 7–8)

Independent of Blocks 1–3 — can parallel if you have time

#	Task	Effort	Status	Original TODO #	Notes
4.1	 GitHub Action running re-fusal_test.py on every push	1 hr	—	#12	One .github/workflows/ci.yml file. Engineering maturity signal.
4.2	 Docker Compose (backend + frontend + MongoDB + Qdrant)	2–3 hrs	—	#9	One command: docker-compose up. Eliminates defense-day “works on my machine” risk.
4.3	 Pre-flight check script (verifies .env, DBs, model before startup)	1 hr	 NEW	—	Script checks: MongoDB reachable? Qdrant loaded? LLM responds? Fail fast with clear error.
4.4	 Query analytics log (append query + category + score to file)	1 hr	 NEW	—	Simple CSV/JSONL log. Shows you think about improvement loops.

Block 4 exit criteria: CI passes, Docker works, pre-flight script runs, analytics log exists.

Sprint Block 5: Report Writing (Days 8–10)

Do these last — they need all the numbers and demos from Blocks 1–4

#	Task	Effort	Status	Original TODO #	Notes
5.1	 Document tempera- ture=0.0 as explicit design choice	30 min	—	#39	“Deterministic outputs reduce eval variance. Intentional trade-off vs. natural phrasing.”

Table 6 – continued

#	Task	Effort	Status	Original TODO #	Notes
5.2	 Document PII/ prompt-injection as out-of-scope	15 min	—	#41	One sentence: “Input sanitization is deliberately deferred for this prototype.”
5.3	 Write “Production Readiness” section	30 min	—	#42	List auth, rate limiting, logging, CI/CD, scaling as known deferred gaps. Shows maturity.
5.4	 Write “Future Work” section	1 hr	 NEW	—	Multi-crop, voice STT/ TTS, domain fine-tuning, real-time prices, farmer feedback loop.
5.5	 Write “Design Choices” section	1 hr	 NEW	—	Why semantic chunking? Why BGE-M3? Why temperature=0.0? Why hard refusal?
5.6	 Data provenance paragraph	30 min	—	#8	Source: Karnataka State Agricultural Department + Kisan Call Centre.

Table 6 – continued

#	Task	Effort	Status	Original TODO #	Notes
5.7	 Integrate all eval tables (retrieval ablation, generation comparison, threshold sweep, refusal accuracy)	2 hrs	 NEW	—	Make tables pretty. Add 95% CI where possible.
5.8	 Final proofread + formatting	2 hrs	 NEW	—	Consistent terminology, no “morphological chunking” (say “semantic chunking”), check all citations.

Block 5 exit criteria: Report is complete, all numbers are fresh, all design choices are justified.

Sprint Block 6: Defense-Day Prep (Day 10+)

Final polish and dry runs

#	Task	Effort	Status	Original TODO #	Notes
6.1	 Dry run the live demo (3–5 preset queries + 1 follow-up)	1 hr	 NEW	—	Practice the flow. Have backup screenshots if live demo fails.

Table 7 – continued

#	Task	Effort	Status	Original TODO #	Notes
6.2	 Prepare backup slides (screenshots of streaming, numeric checker, ablation table)	1 hr	 NEW	—	If live demo crashes, you still have evidence.
6.3	 Prepare 3 “deep dive” talking points	30 min	 NEW	—	Pick 3 things you’ re proud of and can talk for 2 minutes each. Examples: semantic chunking, numeric safety, threshold sweep.
6.4	 Test on a clean machine (friend’ s laptop, or fresh VM)	2 hrs	 NEW	—	Clone repo, follow your own README, verify it works.

Priority Summary (If You Must Cut)

If you only have 1 week, do this subset:

Priority	Task	Why
P0	Test new Groq model + rebuild DB + re-label gold	Foundation —everything else is garbage without this
P0	Query condensation	Biggest demo risk — follow-ups currently fail
P0	Expand eval to n=40	n=16 is not defensible
P1	Streaming	Makes demo feel alive
P1	Run all eval scripts with fresh data	Numbers for report
P1	Adversarial test suite	Shows safety awareness

Table 8 – continued

Priority	Task	Why
P2	GitHub Action	Engineering maturity signal
P2	Docker	Defense-day reliability
P2	Report sections (Design Choices, Future Work, Production Readiness)	Narrative framing
P3	Query expansion	Nice but not critical
P3	Query analytics log	Nice but not critical
P3	Pre-flight script	Nice but not critical

What Got Removed / Deprioritized

These were in your original TODO but are **lower priority** than the new gaps:

Original TODO #	Task	Why Deprioritized
#32	Voice input/output (STT/TTS)	Out of prototype scope. Mention in Future Work.
#33	Multi-crop expansion (ragi, tomato)	Requires new data + re-labeling. Future Work.
#40	Structured extraction for dosages	Track 2 / future work. Regex checker (#5) is sufficient for prototype.

How to Use This File

1. **Copy this into your repo** as `TODO_PIPELINE_ORDERED.md`
2. **Check off items as you complete them** — commit the file with checkmarks
3. **If you get stuck on a task**, mark it with  and move to the next — don't block the sprint
4. **Update the “Last updated” date** at the top when you make changes

This ordering ensures you never do work that gets invalidated by a later task.